


Matmul en NVIDIA GB10: Triton frente a cuBLAS

■ Triton (autotune) ■ cuBLAS



Barras: mediana; bigotes: percentiles 20-80 (do_bench). torch 2.9.0a0+145a3a7bda.nv25.10, triton 3.8.0, CUDA 13.0; job 20423.