


Matmul en NVIDIA GeForce RTX 5090: Triton frente a cuBLAS

--- Pico ficha NVIDIA fp16: 209.5 TFLOP/s Triton (autotune) cuBLAS



Barras: mediana; bigotes: percentiles 20-80 (do_bench). torch 2.9.0a0+145a3a7bda.nv25.10, triton 3.8.0, CUDA 13.0; job 20499.